

Zhuoming Chen

5227 Fifth Avenue #A1, Pittsburgh, PA, United States
(+1) 412-589-5022 chenzhuoming911@gmail.com

EDUCATION

Department of Computer Science, Carnegie Mellon University, PA, USA

Aug 2023-Present

➤ Ph.D. student, Advised by Beidi Chen and Zhihao Jia

Department of Automation, Tsinghua University, Beijing, China

Aug 2019-Jul 2023

➤ B.E. in Automation, Overall GPA: 3.86/4.0

RESEARCH INTERESTS

Machine Learning Systems, Generative AI, Randomized Algorithms

INTERNSHIP EXPERIENCES

FAIR, Meta Platform, Inc | Research Scientist Intern | New York

May 2025-Dec 2025

Mentor: Leon Bottou, Director

Machine Learning Group, Microsoft Research Asia | Research Intern | Beijing

Nov 2021-May 2022

Mentor: Yu Shi, Researcher

PUBLICATIONS

- **Chen, Z.**, Liu, H., Zhou, Y., Zheng, H., & Chen, B. (2026). **Jackpot: Optimal Budgeted Rejection Sampling for Extreme Actor-Policy Mismatch Reinforcement Learning. (ICLR 2026)**
- Agarwal, K., **Chen, Z.**, Luo, C., Chen, Y., Zheng, H., Huang, X., ... & Chen, B. (2026). **MonarchRT: Efficient Attention for Real-Time Video Generation. *arXiv preprint arXiv:2602.12271*.**
- Sadhukhan, R.*, **Chen, Z.***, Zheng, H., Zhou, Y., Strubell, E. and Chen, B., 2025. **Kinetics: Rethinking Test-Time Scaling Laws. (TTODLER25 Best Paper Award, LCFM25 Oral, KDD GenAI25 Oral, NeurIPS 2025)**
- **Chen, Z.**, Sadhukhan, R., Ye, Z., Zhou, Y., Zhang, J., Nolte, N., ... & Chen, B. (2024). **Magicpig: Lsh sampling for efficient llm generation. (ICLR 2025 Spotlight)**
- **Chen, Z.***, May, A.*, Svirschevski, R.*, Huang, Y., Ryabinin, M., Jia, Z. and Chen, B., 2024. **Sequoia: Scalable, Robust, and Hardware-aware Speculative Decoding. (NeurIPS 2024 Spotlight)**
- Svirschevski, R.*, May, A.*, **Chen, Z.***, Chen, B., Jia, Z. and Ryabinin, M. 2024. **SpecExec: Massively Parallel Speculative Decoding for Interactive LLM Inference on Consumer Devices. (NeurIPS 2024)**
- Chen, J., Tiwari, V., Sadhukhan, R., **Chen, Z.**, Shi, J., Yen, I.E.H. and Chen, B., 2024. **MagicDec: Breaking the Latency-Throughput Tradeoff for Long Context Generation with Speculative Decoding. (ICLR 2025)**
- Zhou, Y., **Chen, Z.**, Xu, Z., Lin, V. and Chen, B. 2024. **Sirius: Contextual Sparsity with Correction for Efficient LLMs. (NeurIPS 2024)**
- Luo, C., Zhao, J., **Chen, Z.**, Chen, B. and Anandkumar, A. 2024. **MINI-SEQUENCE TRANSFORMER: Optimizing Intermediate Memory for Long Sequences Training. (NeurIPS 2024)**
- Sun, H., **Chen, Z.**, Yang, X., Tian, Y. and Chen, B., 2024. **TriForce: Lossless Acceleration of Long Sequence Generation with Hierarchical Speculative Decoding. (COLM 2024)**
- Chen, J., **Chen, Z.** and Qian, X., 2023. **GNNPipe: Accelerating Distributed Full-Graph GNN Training with Pipelined Model Parallelism. (HPCA 2025)**
- Miao, X., Oliaro, G., Zhang, Z., Cheng, X., Wang, Z., Wong, R.Y.Y., **Chen, Z.**, Arfeen, D., Abhyankar, R. and Jia, Z., 2023. **SpecInfer: Accelerating Generative LLM Serving with Speculative Inference and Token Tree Verification. (ASPLOS 2024)**
- Shi, Y., Ke, G., **Chen, Z.**, Zheng, S. and Liu, T.Y., 2022. **Quantized training of gradient boosting decision trees. (NeurIPS 2022)**